


Standard Multithreading Support in C++

By: Chris DeVisser

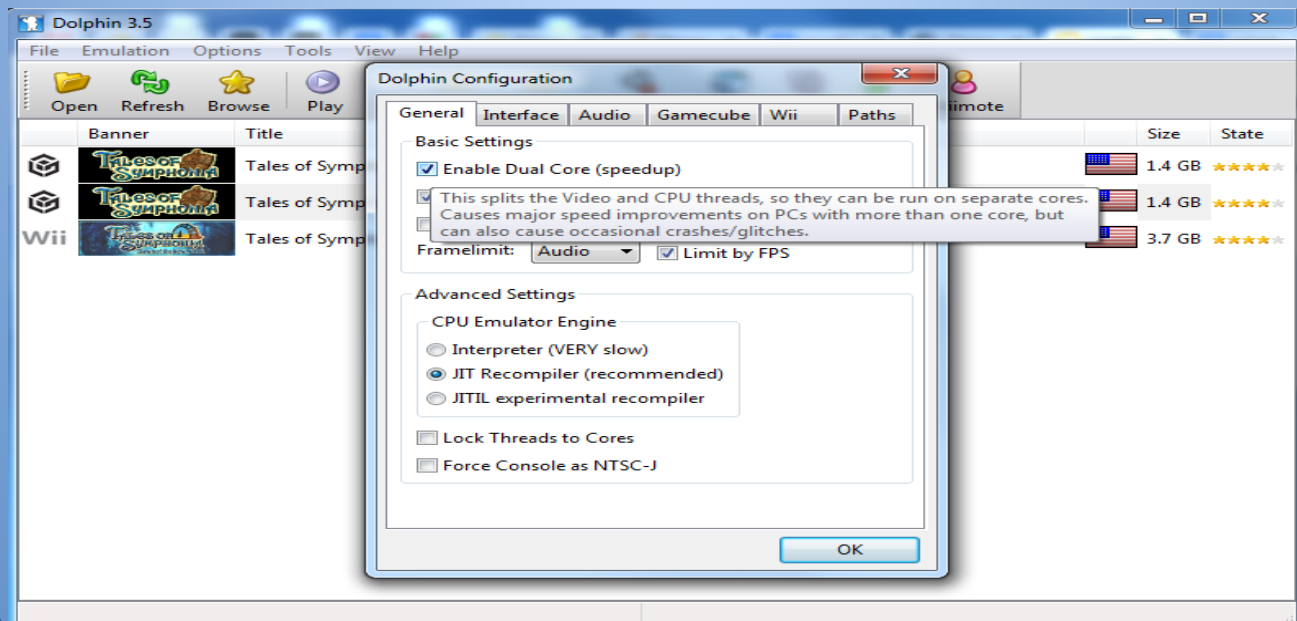


The Importance of Multithreading

- Modern processors have many cores, but often, only one is used in an application
- Significantly speeds up long operations if done correctly
- Increases application responsiveness
- We're moving into a multithreaded world

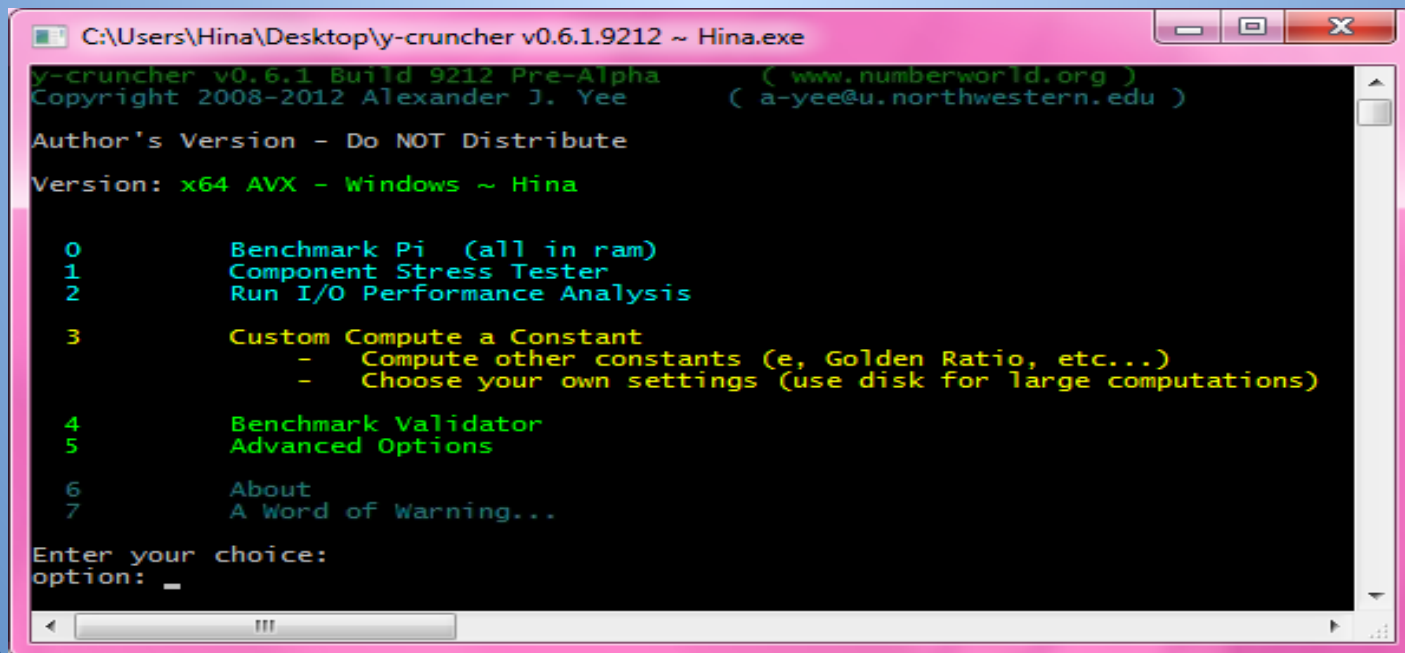
What Uses Multithreading?

- Games
- Number-crunching applications
- Operating Systems
- Most responsive applications



y-cruncher

- Holds the record for most digits of pi calculated, with 10 trillion
- Uses multithreading extensively



```
C:\Users\Hina\Desktop\y-cruncher v0.6.1.9212 ~ Hina.exe
y-cruncher v0.6.1 Build 9212 Pre-Alpha ( www.numberworld.org )
Copyright 2008-2012 Alexander J. Yee ( a-yee@u.northwestern.edu )
Author's Version - Do NOT Distribute
Version: x64 AVX - Windows ~ Hina

0      Benchmark Pi (all in ram)
1      Component Stress Tester
2      Run I/O Performance Analysis

3      Custom Compute a Constant
        - Compute other constants (e, Golden Ratio, etc...)
        - Choose your own settings (use disk for large computations)

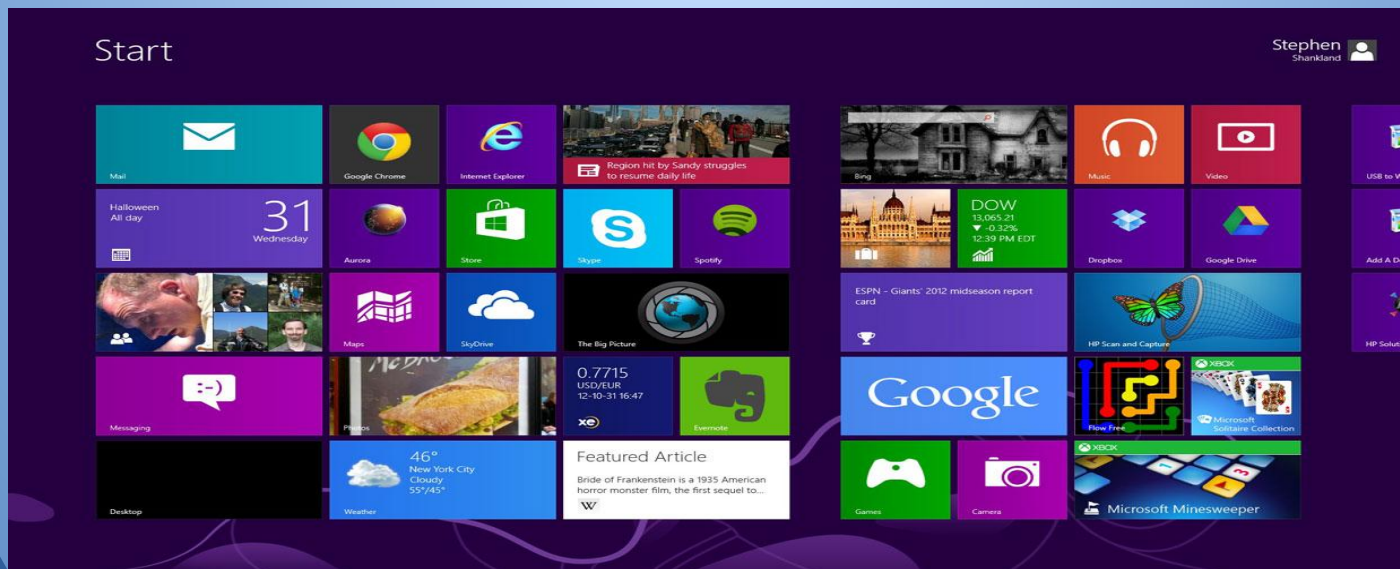
4      Benchmark Validator
5      Advanced Options

6      About
7      A Word of Warning...

Enter your choice:
option: _
```

Windows 8

- "Fast and Fluid" => new policy
- All functions that could take longer than 50ms were made asynchronous
- Apps stay responsive



C++ Concurrency

- C++11 brought standardized concurrency
- Launching another thread is easy:

C++ Concurrency

- C++11 brought standardized concurrency
- Launching another thread is easy:

```
#include <thread>
int main() {
    std::thread t([]{ /*stuff*/ });
    t.join();
}
```

Launching Asynchronous Tasks

- `std::future` holds a task's status

Launching Asynchronous Tasks

- `std::future` holds a task's status

```
auto fut = std::async([]{ /*stuff*/ });  
// concurrent  
fut.wait();
```

Launching Asynchronous Tasks

- `std::future` holds a task's status

```
auto fut = std::async([]{ /*stuff*/ });  
// concurrent  
fut.wait();
```

- Be sure to assign the result of `async`, or the function may not launch a new thread

Preventing Half-Written Data

- Two threads use the same variable

Preventing Half-Written Data

- Two threads use the same variable

```
std::atomic<int> i;
```

Preventing Half-Written Data

- Two threads use the same variable

```
std::atomic<int> i;
```

```
//thread 1  
++i;
```

Preventing Half-Written Data

- Two threads use the same variable

```
std::atomic<int> i;
```

```
//thread 1  
++i;
```

```
//thread 2  
std::cout << i;
```


Tip: Use RAI

- Never forget to clean up again
- Exception-safe

Tip: Use RAII

- Never forget to clean up again
- Exception-safe

```
int i = 0;  
std::mutex m;
```

Tip: Use RAII

- Never forget to clean up again
- Exception-safe

```
int i = 0;  
std::mutex m;
```

```
//threads  
m.lock();  
++i;  
//oops, forgot to unlock
```

Tip: Use RAI

- Never forget to clean up again
- Exception-safe

```
int i = 0;
```

```
std::mutex m;
```

```
//threads
```

```
std::lock_guard<std::mutex> lock(m);
```

```
++i;
```

```
//m unlocked when lock goes out of scope
```

Managing Shared Resource Use

- Two threads use the same resource

Managing Shared Resource Use

- Two threads use the same resource

```
std::ofstream fout("out.txt");  
std::mutex file;
```


Managing Shared Resource Use

- Two threads use the same resource

```
std::ofstream fout("out.txt");  
std::mutex file;
```

```
//thread 1  
std::lock_guard<std::mutex> lock(file);  
fout << "thread 1\n";
```

Managing Shared Resource Use

- Two threads use the same resource

```
std::ofstream fout("out.txt");  
std::mutex file;
```

```
//thread 1  
std::lock_guard<std::mutex> lock(file);  
fout << "thread 1\n";
```

```
//thread 2  
std::lock_guard<std::mutex> lock(file);  
fout << "thread 2\n";
```

Tip: Avoid Deadlocks

- Occur when two threads wait on each other

}

Tip: Avoid Deadlocks

- Occur when two threads wait on each other

//example adapted from <http://pages.cs.wisc.edu/~remzi/Courses/537/Fall2005/Lectures/lecture8.pdf>

```
int X = 0, Y = 0;  
std::mutex x, y;
```

Tip: Avoid Deadlocks

- Occur when two threads wait on each other

//example adapted from <http://pages.cs.wisc.edu/~remzi/Courses/537/Fall2005/Lectures/lecture8.pdf>

```
int X = 0, Y = 0;  
std::mutex x, y;
```

```
//thread 1  
std::lock_guard<std::mutex> lock(x);  
++X;
```

```
//thread 2  
std::lock_guard<std::mutex> lock(y);  
--Y;
```

Tip: Avoid Deadlocks

- Occur when two threads wait on each other

//example adapted from <http://pages.cs.wisc.edu/~remzi/Courses/537/Fall2005/Lectures/lecture8.pdf>

```
int X = 0, Y = 0;  
std::mutex x, y;
```

```
//thread 1  
std::lock_guard<std::mutex> lock(x);  
++X;  
{  
    std::lock_guard<std::mutex> lock(y);  
    ++Y;  
}
```

```
//thread 2  
std::lock_guard<std::mutex> lock(y);  
--Y;  
{  
    std::lock_guard<std::mutex> lock(x);  
    --X;  
}
```


Tip: Avoid Deadlocks

- Occur when two threads wait on each other

//example adapted from <http://pages.cs.wisc.edu/~remzi/Courses/537/Fall2005/Lectures/lecture8.pdf>

```
int X = 0, Y = 0;  
std::mutex x, y;
```

```
//thread 1  
std::lock_guard<std::mutex> lock(x);  
++X;  
{  
    std::lock_guard<std::mutex> lock(y);  
    ++Y;  
}
```

```
//thread 2  
std::lock_guard<std::mutex> lock(y);  
--Y;  
{  
    std::lock_guard<std::mutex> lock(x);  
    --X;  
}
```

Tip: Avoid Deadlocks

```
//thread 1
std::lock_guard<std::mutex> lock(x);
++X;
{
    std::lock_guard<std::mutex> lock(y);
    ++Y;
}
```

```
//thread 2
std::lock_guard<std::mutex> lock(y);
--Y;
{
    std::lock_guard<std::mutex> lock(x);
    --X;
}
```

Tip: Avoid Deadlocks

```
//thread 1
std::lock_guard<std::mutex> lock(x);
++X;
{
    std::lock_guard<std::mutex> lock(y);
    ++Y;
}
```

```
//thread 2
std::lock_guard<std::mutex> lock(y);
--Y;
{
    std::lock_guard<std::mutex> lock(x);
    --X;
}
```

- Imagine this scenario:

Tip: Avoid Deadlocks

```
//thread 1
std::lock_guard<std::mutex> lock(x);
++X;
{
    std::lock_guard<std::mutex> lock(y);
    ++Y;
}

//thread 2
std::lock_guard<std::mutex> lock(y);
--Y;
{
    std::lock_guard<std::mutex> lock(x);
    --X;
}
```

- Imagine this scenario:
 - Thread 1 locks x

Tip: Avoid Deadlocks

```
//thread 1
std::lock_guard<std::mutex> lock(x);
++X;
{
    std::lock_guard<std::mutex> lock(y);
    ++Y;
}

//thread 2
std::lock_guard<std::mutex> lock(y);
--Y;
{
    std::lock_guard<std::mutex> lock(x);
    --X;
}
```

- Imagine this scenario:
 - Thread 1 locks x
 - Thread 2 locks y

Tip: Avoid Deadlocks

```
//thread 1
std::lock_guard<std::mutex> lock(x);
++X;
{
    std::lock_guard<std::mutex> lock(y);
    ++Y;
}

//thread 2
std::lock_guard<std::mutex> lock(y);
--Y;
{
    std::lock_guard<std::mutex> lock(x);
    --X;
}
```

- Imagine this scenario:
 - Thread 1 locks x
 - Thread 2 locks y
 - Thread 1 tries to lock y and waits

Tip: Avoid Deadlocks

```
//thread 1
std::lock_guard<std::mutex> lock(x);
++X;
{
    std::lock_guard<std::mutex> lock(y);
    ++Y;
}
```

```
//thread 2
std::lock_guard<std::mutex> lock(y);
--Y;
{
    std::lock_guard<std::mutex> lock(x);
    --X;
}
```

- Imagine this scenario:
 - Thread 1 locks x
 - Thread 2 locks y
 - Thread 1 tries to lock y and waits
 - Thread 2 tries to lock x and waits

Tip: Avoid Deadlocks

```
//thread 1
std::lock_guard<std::mutex> lock(x);
++X;
{
    std::lock_guard<std::mutex> lock(y);
    ++Y;
}
```

```
//thread 2
std::lock_guard<std::mutex> lock(y);
--Y;
{
    std::lock_guard<std::mutex> lock(x);
    --X;
}
```

- Imagine this scenario:
 - Thread 1 locks x
 - Thread 2 locks y
 - Thread 1 tries to lock y and waits
 - Thread 2 tries to lock x and waits
- Deadlock!

Tip: Avoid Deadlocks

- One solution: lock both

```
int X = 0, Y = 0;  
std::mutex both;
```

```
//thread 1  
std::lock_guard<std::mutex> lock(both);  
++X;  
++Y;
```

```
//thread 2  
std::lock_guard<std::mutex> lock(both);  
--X;  
--Y;
```

Tip: Avoid Deadlocks

- Another solution: use atomic variables

```
std::atomic<int> X = 0, Y = 0;
```

```
//thread 1  
++X;  
++Y;
```

```
//thread 2  
--X;  
--Y;
```

Staying Responsive

- GUIs should always be responsive

Staying Responsive

- GUIs should always be responsive

```
void OnButtonClicked() {  
    tasks_.emplace_back(std::async([] {  
        doLongTask();  
    }));  
}
```

Splitting up Work

- Jobs can be divided

Splitting up Work

- Jobs can be divided

```
std::vector<std::future<void>> futs;  
futs.reserve(numCores);
```

Splitting up Work

- Jobs can be divided

```
std::vector<std::future<void>> futs;  
futs.reserve(numCores);  
auto perThr = totalToCalc / numCores;
```

Splitting up Work

- Jobs can be divided

```
std::vector<std::future<void>> futs;  
futs.reserve(numCores);  
auto perThr = totalToCalc / numCores;  
  
for (int i = 0; i < numCores; ++i) {  
    futs.emplace_back(  
        std::async(calc, perThr*i, (perThr+1) * i)  
    );  
}
```

Compiler Support

- MSVC - Most
- Clang - Full

Compiler Support

- MSVC - Most
- Clang - Full
- GCC
 - Windows - Atomics

Compiler Support

- MSVC - Most
- Clang - Full
- GCC
 - Windows - Atomics
 - Other - Most

C++ Concurrency Awaits
Go have fun with it!

References

- <http://www.numberworld.org/y-cruncher/>
- <http://www.charlespetzold.com/blog/2011/11/Asynchronous-Processing-in-Windows-8.html>
- <http://channel9.msdn.com/Shows/Going+Deep/C-and-Beyond-2012-Herb-Sutter-Concurrency-and-Parallelism>
- <http://en.cppreference.com/w/cpp/thread>
- <http://pages.cs.wisc.edu/~remzi/Classes/537/Fall2005/Lectures/lecture8.pdf>